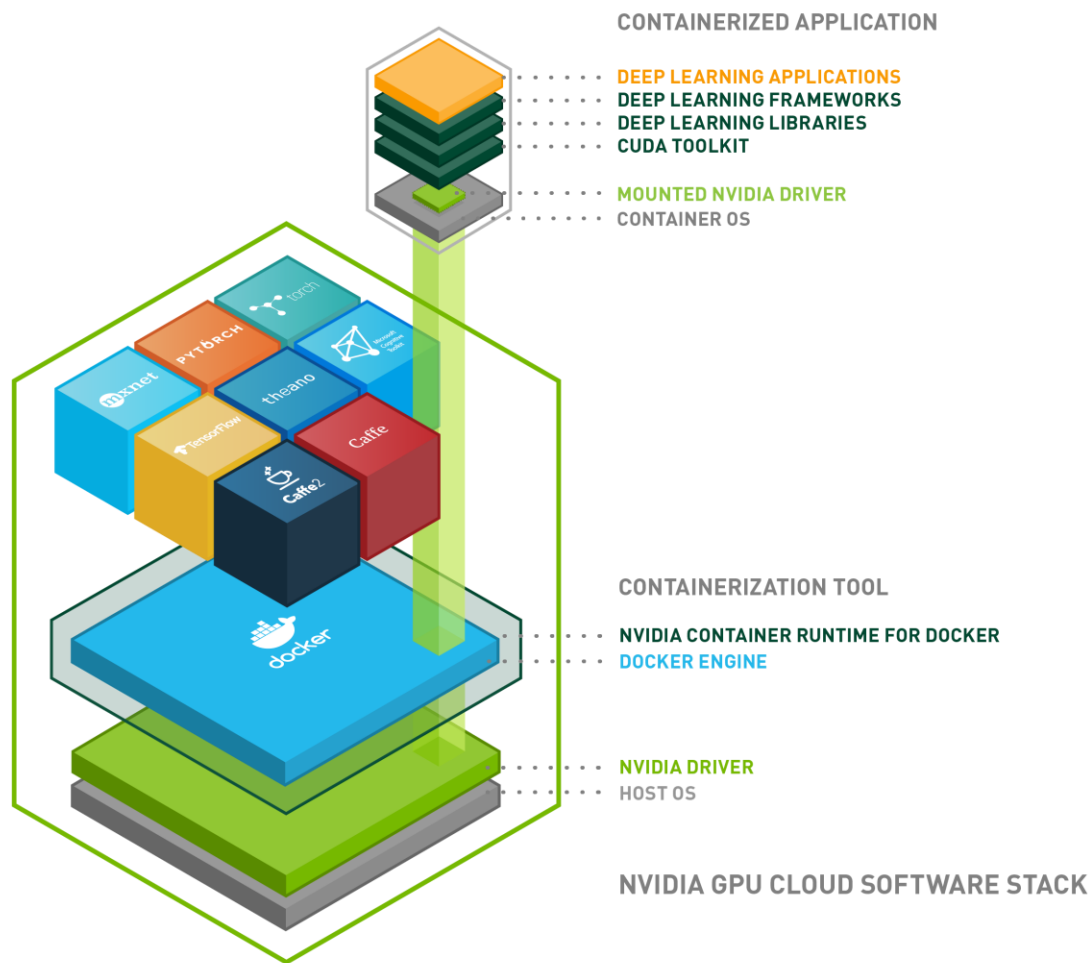




NVIDIA GPU CLOUD

Vishal Mehta



WHY CONTAINERS?

Benefits of Containers:

Simplify deployment of GPU-accelerated software, eliminating time-consuming software integration work

Isolate individual deep learning frameworks and applications


Share, collaborate, and test applications across different environments

CHALLENGES WITH COMPLEX SOFTWARE

Current DIY GPU-accelerated AI and HPC deployments are **complex** and **time consuming** to build, test and maintain

Development of software frameworks by the community is moving **very fast**

Requires high level of **expertise** to manage driver, library, framework dependencies



NVIDIA GPU CLOUD

Simple access to a comprehensive catalog of GPU-accelerated software

Discover 30 GPU-Accelerated Containers

Deep learning, third-party managed HPC applications, NVIDIA HPC visualization tools, and partner applications

Innovate in Minutes, Not Weeks

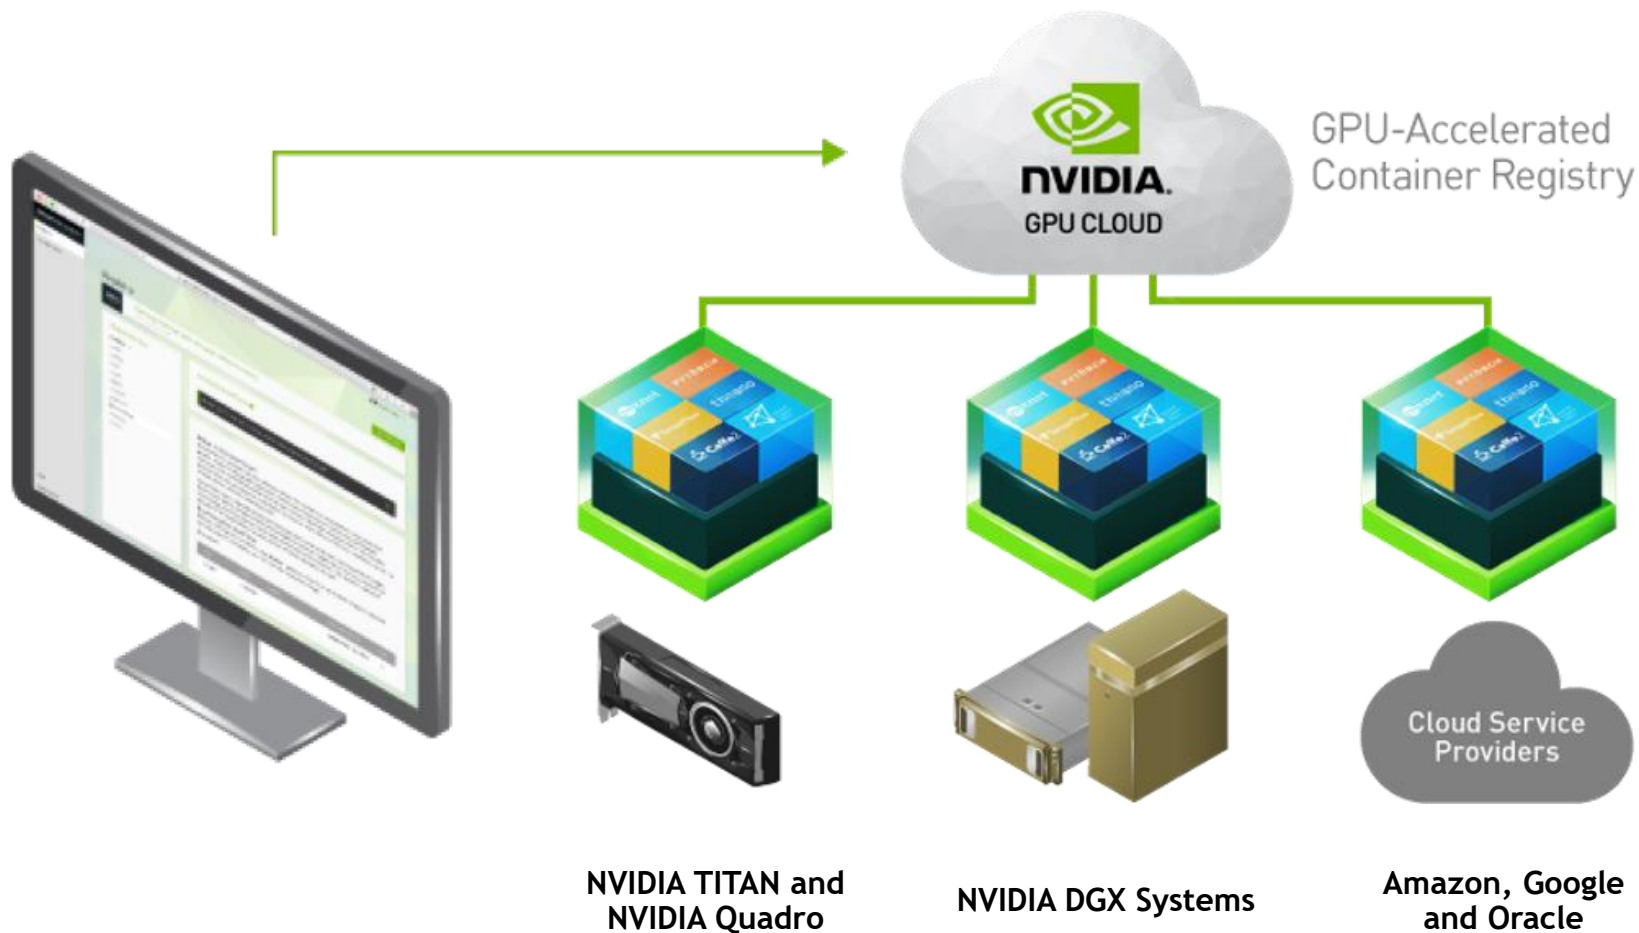
Get up and running quickly and reduce complexity

Access from Anywhere

Use on PCs with NVIDIA Volta or Pascal™ architecture GPUs, NVIDIA DGX Systems, and supported cloud providers



FROM DESKTOP, TO DATA CENTER, TO CLOUD



NGC GPU-OPTIMIZED DEEP LEARNING CONTAINERS

A comprehensive catalog of deep learning software

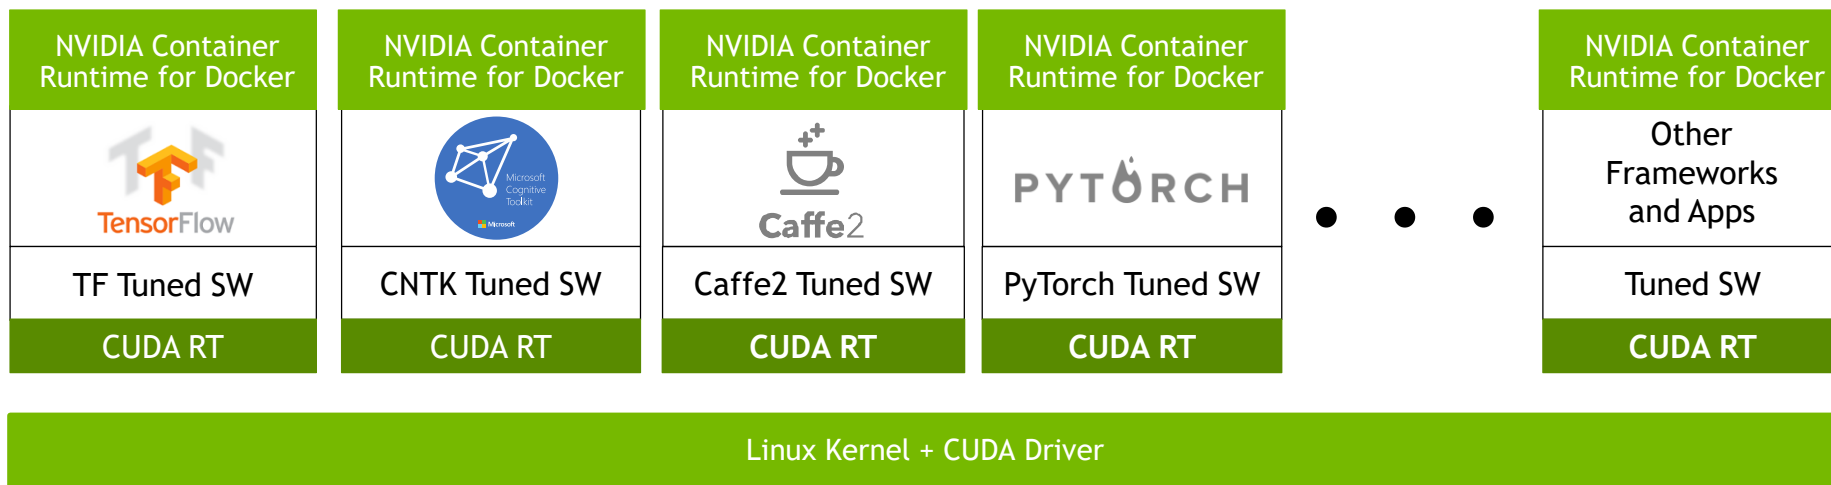
- ▶ NVCaffe
- ▶ Caffe2
- ▶ Microsoft Cognitive Toolkit (CNTK)
- ▶ DIGITS
- ▶ MXNet
- ▶ PyTorch
- ▶ TensorFlow
- ▶ Theano
- ▶ Torch
- ▶ CUDA (base level container for developers)
- ▶ NVIDIA TensorRT inference accelerator with ONNX support



ALWAYS UP-TO-DATE

Monthly updates from NVIDIA to deep learning containers

Containerized Applications



END-TO-END PRODUCT FAMILY

TRAINING

DESKTOP



GPU-Accelerated
Container Registry



TITAN V



DGX Station

DATA CENTER



GPU-Accelerated
Container Registry




DGX-1 Server



TESLA V100

INFERENCE

DATA CENTER



TESLA P4



TESLA V100

EMBEDDED



JETPACK SDK



Jetson

AUTOMOTIVE



DriveWorks SDK



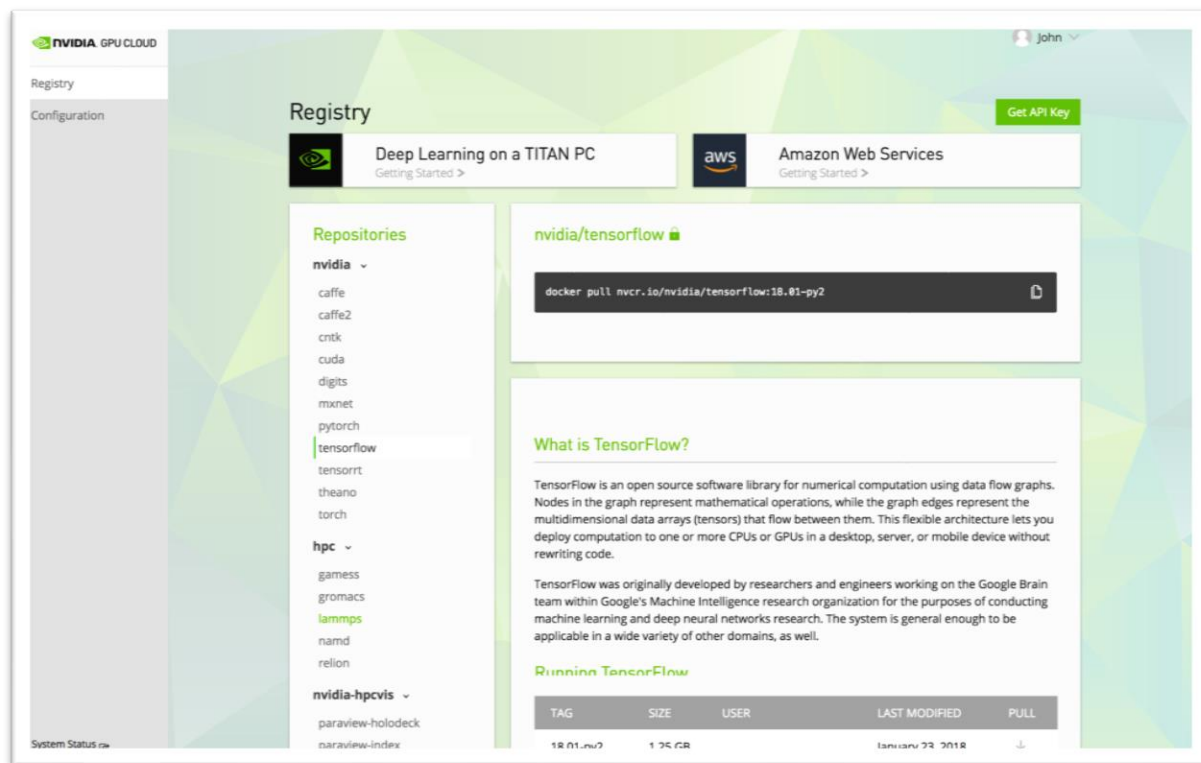
Drive PX

GET STARTED TODAY WITH NGC

Sign up for no cost access

To learn more about all of the GPU-accelerated software on NVIDIA GPU Cloud, visit:
nvidia.com/cloud

To sign up, go to:
nvidia.com/ngcsignup



NGC CONTAINERS ON CSCS

In just 1,2,3.....

CSCS uses shifter to manage containers. Shifter essentially does not allow root access and maintains user privileges.

Run your first NGC Tensorflow container in 3 simple steps

- 1) `shifter pull --login nvcr.io/nvidia/tensorflow:18.03-py3`
- 2) Username: `$oauthtoken` Password: `<API_KEY>`
- 3) `srun -N1 -C gpu --pty shifter run --mount=type=bind,source=$HOME,destination=$HOME nvcr.io/nvidia/tensorflow:18.03-py3 bash`



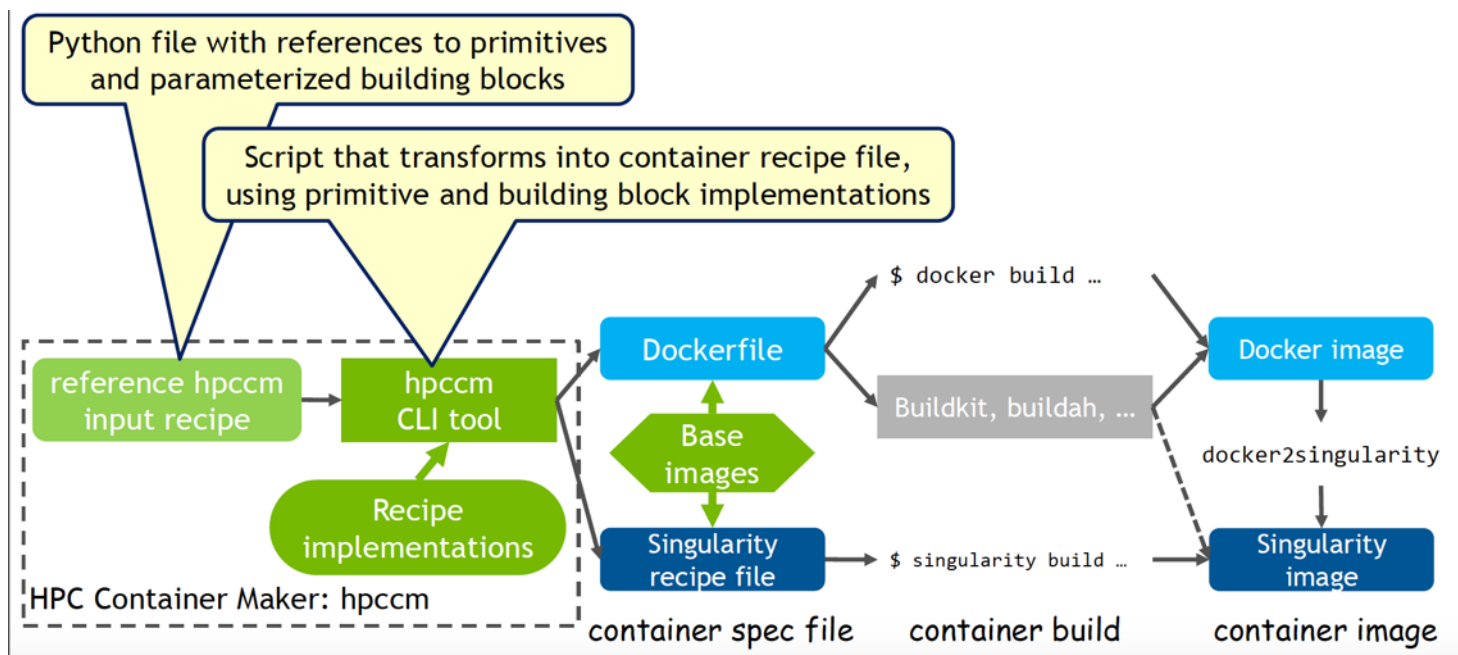
HPC CONTAINER MAKER

Vishal Mehta

HPC CONTAINER MAKER - HPCCM

“h-p-see-um”

- HPC Container Maker (HPCCM) generates container specification files (Dockerfiles or Singularity recipe) based on recipes
- A recipe specifies the series of steps to be performed when building a container



HPC CONTAINER MAKER - HPCCM

- Container implementation abstraction
 - The same recipe file generates specification files for Docker or Singularity
- Availability of full programming languages
 - A recipe is Python code. This means that you can use the full power of Python in a recipe for conditional branching, input validation, searching the web for the latest version of a component, etc.
- Higher level abstraction
 - Provides building blocks to simplify recipes and encapsulate best practices
- Container Maker generates human readable Dockerfiles and Singularity recipe files

FEATURES

BASE IMAGE & PACKAGES

- Ubuntu 16.04
- Cent OS

COMPILER

- PGI COMMUNITY EDITION
- GNU
- NVCC (CUDA TOOLKIT)

SOFTWARE - MORE ON REQUEST

- CMAKE
- HDF5
- NETCDF
- MLNX OFED
- OpenMPI
- MVAPICH2
- Python
- FFTW
- MKL

FULL SCALE APPLICATIONS

- MILC
- GROMACS

```
# Choose a base image

Stage0.baseimage('ubuntu:16.04')

# Install GNU Compilers

Stage0 += apt_get(ospackages=['gcc', 'g++', 'gfortran'])
```

hpccm.py --recipe recipes/basic.py --format docker

```
FROM ubuntu:16.04

RUN apt-get update -y && \
apt-get install -y --no-install-recommends \
gcc \
g++ \
gfortran && \
rm -rf /var/lib/apt/lists/*
```

hpccm.py --recipe recipes/basic.py --format singularity

```
Bootstrap: docker

From: ubuntu:16.04
%post
apt-get update -y
apt-get install -y --no-install-recommends \
gcc \
g++ \
gfortran
rm -rf /var/lib/apt/lists/*
```

AVAILABILITY OF A FULL PROGRAMMING LANGUAGE

- Full power of Python in a recipe for conditional branching, input validation, searching the web for the latest version of a component, etc.

For example, the LAMMPS application may be built in single, double, or mixed precision mode. (hpccm.py --userarg LAMMPS_PRECISION=...)

```
# get and validate precision
VALID_PRECISION = ['single', 'double', 'mixed']
precision = USERARG.get('LAMMPS_PRECISION', 'single')
if precision not in VALID_PRECISION:
    raise ValueError('Invalid precision')
...

Stage0 += shell(commands=['make -f Makefile.linux.{0}'.format(precision), ...])

...
```

HIGHER LEVEL ABSTRACTION

Building blocks to simplify recipes and encapsulate best practices

```
Stage0 += openmpi(cuda=True, infiniband=True,  
  
               prefix='/usr/local/openmpi', version='3.0.0')
```

```
# OpenMPI version 3.0.0 RUN  
apt-get update -y && \  
apt-get install -y --no-install-recommends \  
    file \  
    hwloc \  
    openssh-client \  
    wget && \  
rm -rf /var/lib/apt/lists/*  
  
RUN mkdir -p /tmp && wget -q --no-check-certificate -P /tmp https://www.open-  
mpi.org/software/ompi/v3.0/downloads/openmpi-3.0.0.tar.bz2 && \  
    tar -x -f /tmp/openmpi-3.0.0.tar.bz2 -C /tmp -j && \  
    cd /tmp/openmpi-3.0.0 && ./configure --prefix=/usr/local/openmpi --disable-  
getpwuid --enable-orterun-prefix-by-default --with-cuda --with-verbs && \  
    make -j4 && \  
    make -j4 install && \  
    rm -rf /tmp/openmpi-3.0.0.tar.bz2 /tmp/openmpi-3.0.0  
  
ENV PATH=/usr/local/openmpi/bin:$PATH \  
LD_LIBRARY_PATH=/usr/local/openmpi/lib:$LD_LIBRARY_PATH
```

```
# OpenMPI version 3.0.0  
%post  
    apt-get update -y  
    apt-get install -y --no-install-recommends \  
        file \  
        hwloc \  
        openssh-client \  
        wget  
    rm -rf /var/lib/apt/lists/*  
  
%post  
    mkdir -p /tmp && wget -q --no-check-certificate -P /tmp https://www.open-  
mpi.org/software/ompi/v3.0/downloads/openmpi-3.0.0.tar.bz2  
    tar -x -f /tmp/openmpi-3.0.0.tar.bz2 -C /tmp -j  
    cd /tmp/openmpi-3.0.0 && ./configure --prefix=/usr/local/openmpi --disable-getpwuid  
--enable-orterun-prefix-by-default --with-cuda --with-verbs  
    make -j4  
    make -j4 install  
    rm -rf /tmp/openmpi-3.0.0.tar.bz2 /tmp/openmpi-3.0.0  
  
%environment  
    export PATH=/usr/local/openmpi/bin:$PATH  
    export LD_LIBRARY_PATH=/usr/local/openmpi/lib:$LD_LIBRARY_PATH
```

TOOLCHAIN SUPPORT

Using toolchain for building blocks

```
GNU = gnu()  
Stage0 += GNU  
TC = GNU.toolchain  
TC.CUDA_HOME = /usr/local/cuda  
  
OMPI = openmpi(version='3.0.0',toolchain=TC)  
Stage0 += OMPI  
  
FFTW = fftw(version='3.3.7',toolchain=TC)  
Stage0 += FFTW  
  
HDF5 = hdf5(version='1.10.1',toolchain=TC)  
Stage0 += HDF5  
  
NETCDF = netcdf(version='4.6.1',toolchain=TC)  
Stage0 += NETCDF
```


USING HPCCM FOR CSCS

ON LOCAL LINUX BOX:

- 1) Git clone <https://github.com/NVIDIA/hpc-container-maker>
- 2) `./hpccm.py --format docker --recipe hpcbase-gnu-mvapich2.py > Dockerfile`
- 3) `Docker build -t container - < Dockerfile`
- 4) `Docker save container -o container.tar.gz`

ON CSCS: (https://user.cscs.ch/user_services/containers/advanced_shifter/)

- 1) `Srun -N 1 -C gpu shifter load container.tar.gz container`
- 2) `srun -N1 -C gpu --pty shifter run --mount=type=bind,source=$HOME,destination=$HOME
load/library/container:latest bash`

